

Name: Saurabh Bansode  
<[sbansod@ncsu.edu](mailto:sbansod@ncsu.edu)>  
Unity ID: 200539958

## ECE 566 Project 2: Common Subexpression Elimination Plus Simple Load/Store Optimization report

**Q1. Describe your implementation and any embellishments you made over my description.**

Function name	Description
<b>DeadInstRemoval(M)</b>	<p>In each basic block, code iterates over the instructions.</p> <p>Checks for dead instruction, if dead removes from parent, if not, simplifies the instruction using <b>simplifyInstruction</b> API.</p> <p>Simplified instruction is then substituted for the older instruction using <b>replaceAllUsesWith</b> API</p>
<b>local_CSE(M)</b>	<p>While iterating over all instructions, check if each instruction, lets call it <b>INST</b> is any one of the below and sets a boolean flag: 1. Load 2. Store 3. Alloca 4. Call 5. Return 6. Branch</p> <p>If the flag is set, it does not perform any CSE.</p> <p>If the flag is not set, then all the instructions that are after <b>INST</b> are checked for type as load, branch, store etc is checked.</p> <p>Comparison done using <b>isIdenticalTo</b> API.</p> <p>If identical, the matched instruction that is later in program order and is inserted in a vector data structure. All uses of the later instruction are replaced with <b>INST</b>.</p> <p>Once all the instructions within a basicblock are checked, the vector is deleted. For each instruction deleted from the vector with the help of <b>eraseFromParent</b>, <b>CSEElim</b> is incremented once.</p> <p>The instruction <b>INST</b> and the basic block is then passed to <b>global_CSE</b>, explained in the next step</p>
<b>global_CSE(inst_in, basicblock_in)</b>	<p>By utilising <b>getParent</b>, and <b>getDescendants</b> API, I successfully created the Dominator Tree for each function the basic block <b>basicblock_in</b> belonged to.</p>

	<p>Then I checked if the instruction received as argument, <b>inst_in</b>, was present in any other basic blocks of the function <b>basicblock_in</b> was a part of using the <code>isIdenticalTo</code> API. If any hits, I replaced the matched instruction with <b>inst_in</b>, using <b>replaceAllUsesWith</b> API. I inserted the instruction in a vector. Incremented the <b>CSEElim</b> counter for each instruction pushed onto the stack.</p> <p>At the end of iteration for each basic block, I deleted the duplicate instructions using <b>eraseFromParent</b> API.</p> <p>The combination of localCSE and globalCSE made sure that instructions repeating both locally and globally are replaced.</p>
<b>elim_red_load(M)</b>	<p>Iterated over each instruction of all basic blocks in the module. For each instruction being iterated, let's call it <b>INST</b>, I <b>checked if it was a load</b>.</p> <p>If yes, I iterated over remaining instructions in that basic block. If the instruction being iterated over was a <b>store or call instruction, I broke out of the for loop</b> and looked for another instruction altogether.</p> <p>If the instruction being iterated over was a load, and it was a <b>non-volatile</b> instruction, I obtained the type of the instructions using <b>getType()</b> API and compared them.</p> <p>If all 3 conditions matched, I checked if the 0th operand of the load instruction, which is the address, are the same, which means it's a redundant load.</p> <p>I replaced the later load with an earlier load. And pushed the later load in a vector to be deleted later. Incremented the <b>CSELdElim</b> counter every time an instruction was pushed onto the vector.</p>
<b>elim_red_stores(M)</b>	<p>When iterating over all instructions of the basic block, check if the instruction is a store. If yes type-casted it to a store instruction using <b>dyn_cast</b> and named it <b>s_inst</b>.</p> <p>Used a for loop to iterate over the <b>remaining instructions</b> (lets call <b>r_inst</b>) in the block to see if any redundant stores/loads are present.</p> <p>Checked if</p> <ol style="list-style-type: none"> <li>1. <b>r_inst</b> and <b>s_inst</b> are have same address</li> <li>2. <b>s_inst</b>'s operand type = type of <b>r_inst</b></li> <li>3. <b>r_inst</b> is non-volatile</li> </ol>

	<p>4. <b>r_inst</b> is a load instruction If all conditions match, increment <b>CSEStore2Load</b>, replace all uses of r_inst with s_inst using <b>replaceAllUsesWith</b> API.</p> <p>Second condition was checked for any store elimination</p>
--	--

**Q2. Collect data per benchmark that compares the number of instructions, the total number of loads, the total number of stores, the counters you collect in the CSE pass**

**Benchmark data focusing on counters, with just CSE enabled and no mem2reg**

Bench- mark Name	Instructi ons	Loa ds	Stor es	CSED ead	CSEE lim	CSE Simp lify	CSELd Elim	CSEStore2 Load	CSESt Elim
adpcm	418	122	81	1	0	4	0	0	0
arm	744	123	116	0	14	40	24	0	0
basicma th	566	156	100	2	10	19	11	0	0
bh	3104	825	494	51	26	83	67	0	0
bitcount	637	158	98	1	0	7	17	0	0
crc32	32141	33	29	0	0	0	3	1	0
dijkstra	319	94	51	0	0	8	1	0	0
em3d	1221	418	192	1	0	35	1	0	0
fft	699	208	102	1	16	8	8	0	0
hanoi	96	29	16	0	0	1	0	0	0
hello	4	0	1	0	0	0	0	0	0
kmp	529	146	71	0	0	16	27	0	0
l2lat	89	24	15	4	0	1	1	0	0
patricia	1043	370	108	1	4	46	8	0	0
qsort	142	36	16	1	0	3	2	0	0
sha	624	186	99	0	5	36	24	1	0

smatrix	291	75	31	2	0	1	22	0	0
sql	171601	570 82	218 97	300	414	3960	1834	0	0
susan	12181	420 9	143 8	5	68	22	364	0	0

### Timing without mem2reg

adpcm	2.32
arm	0
basicmath	0.08
bh	1.11
bitcount	0.23
crc	0
dijkstra	0.05
em3d	0.4
fft	0.07
hanoi	5.53
hello	0
kmp	0.21
l2lat	0.03
patricia	0.11
qsort	0.04
sha	0
smatrix	7.35
sql	0
susan	0.7

### Benchmark data focusing on counters, with just CSE enabled AND mem2reg

Bench- mark Name	Instructi ons	Loa ds	Stor es	CSED ead	CSEE lim	CSE Simp lify	CSELd Elim	CSEStore2 Load	CSESt Elim
adpcm	240	15	7	2	7	6	0	0	0
arm	394	47	18	0	33	42	2	0	0
basicma th	326	24	12	1	20	19	1	0	0
bh	1832	192	142	1	168	84	5	0	0

bitcount	426	51	18	1	7	7	0	0	0
crc32	83	8	4	0	0	0	0	0	0
dijkstra	225	48	24	0	7	8	1	0	0
em3d	634	117	43	3	50	36	0	0	0
fft	385	38	24	0	52	9	0	0	0
hanoi	52	6	4	0	1	1	0	0	0
hello	2	0	0	0	0	0	0	0	0
kmp	331	47	20	0	38	16	11	0	0
l2lat	57	8	1	0	1	1	0	0	0
patricia	638	133	30	0	76	48	4	0	0
qsort	89	13	4	1	10	3	0	0	0
sha	371	41	28	0	46	36	1	0	0
smatrix	210	34	10	0	20	1	5	0	0
sql	103159	163 40	584 3	188	7076	4087	239	5	0
susan	6544	102 1	157	0	1210	34	34	0	0

### Timing with mem2reg

adpcm	2.89
arm	0
basicmath	0.09
bh	1.31
bitcount	0.28
crc	0.39
dijkstra	0.09
em3d	0.45
fft	0.07
hanoi	4.04
hello	0
kmp	0.23
l2lat	0.05

patricia	0.14
qsort	0.05
sha	0.06
smatrix	7.07
sql	0
susan	1.59

### Q3. Explain the difference in results for these two configurations using your data for Q2


Difference in results for the configs used in Q2 is below:

Just CSE, NO mem2reg	CSE and mem2reg
More Instructions, more stores, more loads	Compared to when we do not apply memory to register promotion, we get very few instructions, very few stores and loads. When we apply mem2reg optimization
Looking at susan and emd3 benchmark, I conclude that without mem2reg optimization, there are more dead instruction removal	With memory to register optimization applied, there are more CSE eliminations for susan, fft and bh benchmarks
There is very little information on store elimination.	Store elimination is practically not present in this data

### Q4. Compare the output counters you collected.


**CSEDead:** Apart from susan and bitscount, there is not a significant reduction in dead instructions.

When Mem2reg is applied along with CSE, we can see lot less instructions being removed.




**CSEElim:** For sql benchmark, when mem2reg is applied, there is humongous amount of instructions


Being eliminated as means for CSEElim. Other benchmarks have minimal differences.



**CSESimplify:** For sql benchmark, there is significant increase in instructions being simplified when both CSE and memory 2 register optimizations are applied.




For **CSEldElim**, bh and sql benchmarks display that applying memory 2 register optimization along with CSE improves in reducing the number of redundant load instructions.



**CSEStore2Load** – For CSE alone, only crc32 and sha has any instructions removed.

Whereas when CSE+mem2reg is applied, both crc32 and sha show no instructions being removed, however sql shows huge improvement.





**CSEStElim:** For CSE Store Eliminations, neither CSE alone nor CSE+mem2reg shows any instructions being removed. This means that my implementation for CSE store elimination is flawed.

